

# Lecture 10

R and the tidyverse // regression

---

Ivan Rudik  
AEM 4510

# Roadmap

- What is R?
- What is the tidyverse?
- How do we import and manipulate data?

Our goal is to take a hands on approach to learning how we do environmental economics research

A good chunk of this lecture comes from Grant McDermott's [data science for economists notes](#), and [RStudio education](#)

# RStudio Cloud

---

# Getting started

We will be using [rstudio.cloud](#) for our coding

# Getting started

We will be using [rstudio.cloud](#) for our coding

Why?

# Getting started

We will be using **rstudio.cloud** for our coding

Why?

You don't need to download/install anything

# Getting started

We will be using **rstudio.cloud** for our coding

Why?

You don't need to download/install anything

I can prepare the packages and code and make it easy to download

# Getting started

We will be using [rstudio.cloud](#) for our coding

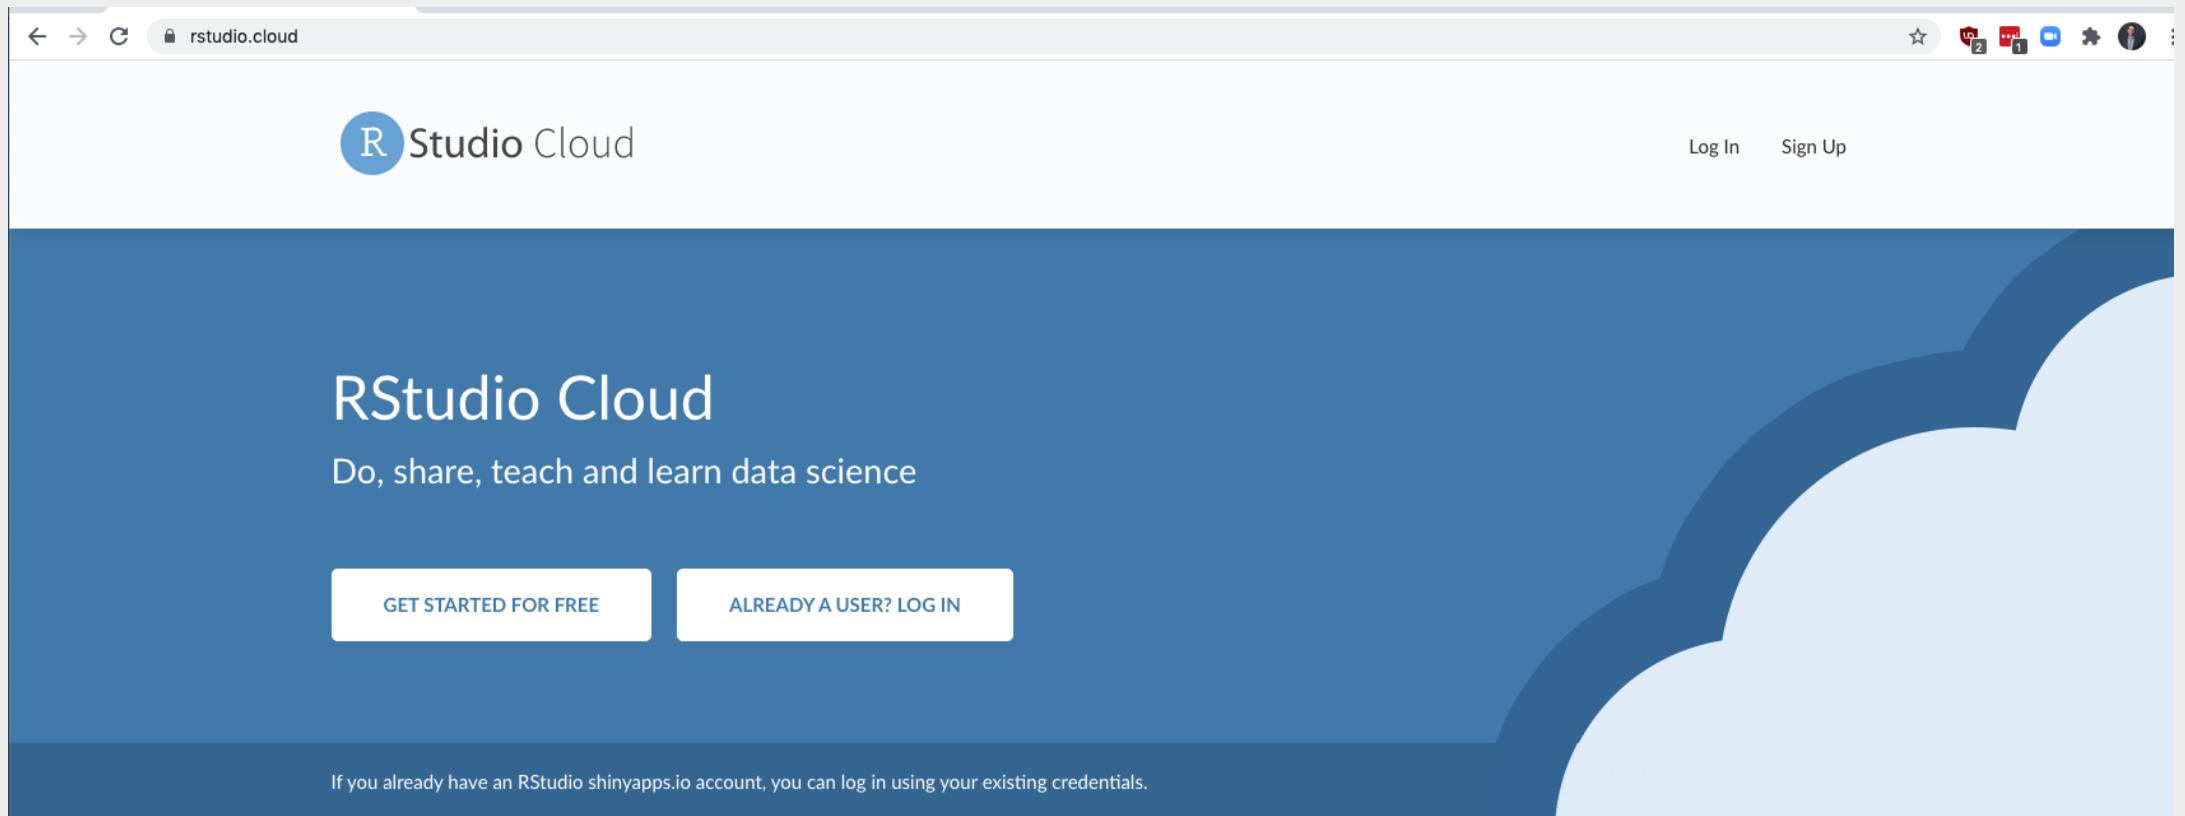
Why?

You don't need to download/install anything

I can prepare the packages and code and make it easy to download

Let's get everything going...

# Getting started: login



The screenshot shows the RStudio Cloud login page. At the top, there's a header bar with browser controls, a lock icon, and the URL "rstudio.cloud". To the right of the URL are icons for a star, a user profile, and other settings. Below the header, the RStudio Cloud logo is on the left, and "Log In" and "Sign Up" buttons are on the right. The main section has a blue background with a white cloud graphic on the right. It features the text "RStudio Cloud" and "Do, share, teach and learn data science". Below this are two buttons: "GET STARTED FOR FREE" and "ALREADY A USER? LOG IN". A note at the bottom of this section says, "If you already have an RStudio shinyapps.io account, you can log in using your existing credentials." The bottom part of the page has a white background with a dark blue footer bar.

R Studio Cloud

Log In    Sign Up

# RStudio Cloud

Do, share, teach and learn data science

GET STARTED FOR FREE    ALREADY A USER? LOG IN

If you already have an RStudio shinyapps.io account, you can log in using your existing credentials.

## Data science without the hardware hassles

RStudio Cloud is a lightweight, cloud-based solution that allows anyone to do, share, teach and learn data science online.

- Analyze your data using the RStudio IDE, directly from your browser.

---

\$ AVAILABLE PRICING PLANS

---

🕒 RSTUDIO CLOUD GUIDE

---

🌐 RSTUDIO.COM

---

# Getting started: new project from github

The screenshot shows the RStudio Cloud interface at [rstudio.cloud/projects](https://rstudio.cloud/projects). The left sidebar includes sections for Spaces (Your Workspace selected), Learn (Guide, What's New, Primers, Cheat Sheets), Help (Current System Status, RStudio Community), and Info (Plans & Pricing, Terms and Conditions). The main content area displays 'Your Projects' with a 'Sort By name' dropdown and a 'no projects' message. A 'New Project' button and a 'New Project from Git Repository' link are available. On the right, an 'Info' panel for Ivan Rudik states it's a personal workspace, provides a guide link, and shows account usage: 0 projects and 0.4 hours out of 15. It also offers an upgrade link.

Your Workspace Projects About

R Studio Cloud Ivan Rudik

Spaces

Your Workspace New Space

Learn

Guide What's New Primers Cheat Sheets

Help

Current System Status RStudio Community

Info

Plans & Pricing Terms and Conditions

Your Projects

Sort By name

New Project

New Project from Git Repository

no projects

This is your personal workspace.

Learn more about [Your Workspace](#) in the [Guide](#).

Account Usage

Projects: 0 of 15

Project hours: 0.4 of 15

[Upgrade your account](#) to use more project hours, create more projects, and access other premium features.

6 / 156

# Getting started: new project from github

The screenshot shows the RStudio Cloud interface at [rstudio.cloud/projects](https://rstudio.cloud/projects). The left sidebar includes links for 'Your Workspace' (selected), 'New Space', 'Guide', 'What's New', 'Primer', 'Cheat Sheets', 'Current System Status', and 'RStudio Community'. The main area displays 'Your Projects' with a sorting dropdown set to 'By name'. A central modal window titled 'New Project from Git Repository' contains a text input field with the URL <https://github.com/irudik/aem6510>. Below the input is an 'OK' button. To the right of the modal is an 'Info' panel with the message: 'This is your personal workspace. Learn more about Your Workspace in the Guide.' At the bottom right of the page is an 'Account Usage' section showing 'Projects: 0 of 15' and 'Project hours: 0.4 of 15', along with a link to upgrade the account.

# Getting started: wait for deployment

The screenshot shows the RStudio Cloud interface with the URL `rstudio.cloud/project/1795327` in the address bar. The left sidebar is visible, showing sections for Spaces, Learn, and Help. The main content area displays a large, semi-transparent circular overlay with the text "Deploying Project" and a progress bar at the bottom.

R Studio Cloud

Your Workspace / aem6510

Spaces

- Your Workspace
- New Space

Learn

- Guide
- What's New
- Primers
- Cheat Sheets

Help

- Current System Status
- RStudio Community

Info

- Plans & Pricing
- Terms and Conditions

Deploying Project

8 / 156

# Click on class-code in bottom-right

The screenshot shows the RStudio Cloud interface for a workspace named 'aem6510'. The left sidebar contains links for 'Your Workspace', 'New Space', 'Guide', 'What's New', 'Primers', 'Cheat Sheets', 'Help', 'Current System Status', and 'RStudio Community'. The main area has tabs for 'Console' (active), 'Terminal', and 'Jobs'. The 'Console' tab displays the R startup message and a single command prompt line starting with '>'. The 'Environment' tab shows an 'Empty Environment'. The 'Files' tab displays a file tree with the following contents:

| Name          | Size    | Modified              |
|---------------|---------|-----------------------|
| ..            |         | Oct 20, 2020, 1:48 PM |
| .gitignore    | 570 B   | Oct 20, 2020, 1:48 PM |
| .Rhistory     | 0 B     | Oct 20, 2020, 1:48 PM |
| aem6510.Rproj | 205 B   | Oct 20, 2020, 1:48 PM |
| class-code    |         |                       |
| lecture-notes |         |                       |
| LICENSE       | 1 KB    | Oct 20, 2020, 1:48 PM |
| README.md     | 11.9 KB | Oct 20, 2020, 1:48 PM |

# Click on class-code.Rproj

The screenshot shows the RStudio Cloud interface for a workspace named 'aem6510'. The left sidebar contains links for 'Spaces', 'Learn', 'Help', and 'Info'. The main area has tabs for 'Console', 'Terminal', and 'Jobs'. The 'Console' tab shows the message 'Restarting R session...'. The 'Environment' tab shows 'Global Environment' with the message 'Environment is empty'. The 'Files' tab displays a file tree for 'Cloud > project > class-code' with the following contents:

| Name             | Size    | Modified              |
|------------------|---------|-----------------------|
| ..               |         | Oct 20, 2020, 1:48 PM |
| .Rprofile        | 26 B    | Oct 20, 2020, 1:48 PM |
| class-code.Rproj | 205 B   | Oct 20, 2020, 1:48 PM |
| code-here.R      | 101 B   | Oct 20, 2020, 1:48 PM |
| renv             |         |                       |
| renv.lock        | 17.6 KB | Oct 20, 2020, 1:48 PM |

# Click yes

The screenshot shows the RStudio Cloud interface. On the left is a sidebar with links for Spaces, Learn, Help, and Info. The main area has tabs for Console, Terminal, and Jobs. The Console tab shows the message "Restarting R session...". To the right are tabs for Environment, History, Connections, Git, and Tutorial. The Environment tab displays "Global Environment" with the message "Environment is empty". Below these are tabs for Files, Plots, Packages, Help, and Viewer. A file browser window is open, showing a project structure with files like ".Rproj", ".R", and "renv.lock". A "Confirm Open Project" dialog box is overlaid on the interface, asking "Do you want to open the project /cloud/project/class-code?". It has "Yes" and "No" buttons. The status bar at the bottom right shows "11 / 156".

rstudio.cloud/project/1795327

R Studio Cloud Your Workspace / aem6510

File Edit Code View Plots Session Build Debug Profile Tools Help

Spaces

Your Workspace

New Space

Learn

Guide

What's New

Primers

Cheat Sheets

Help

Current System Status

RStudio Community

Info

Plans & Pricing

Terms and Conditions

Console Terminal Jobs

/cloud/project/

Restarting R session...

Environment History Connections Git Tutorial

Import Dataset

Global Environment

Environment is empty

Files Plots Packages Help Viewer

New Folder Upload Delete Rename More

Cloud > project > class-code

Confirm Open Project

Do you want to open the project /cloud/project/class-code?

Yes No

|           | Size    | Modified              |
|-----------|---------|-----------------------|
| .Rproj    | 26 B    | Oct 20, 2020, 1:48 PM |
| .R        | 205 B   | Oct 20, 2020, 1:48 PM |
| renv.lock | 101 B   | Oct 20, 2020, 1:48 PM |
|           | 17.6 KB | Oct 20, 2020, 1:48 PM |

11 / 156

# Check package status (not required)

The screenshot shows the RStudio Cloud interface for a workspace named "aem6510".

**Console Tab:**

```
R version 4.0.3 (2020-10-10) -- "Bunny-Wunnies Freak Out"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-pc-linux-gnu (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

* Downloading renv 0.12.0 from CRAN ... trying URL 'http://package-proxy/src/contrib/renv_0.12.0.tar.gz'
Content type 'application/x-tar' length 1011356 bytes (987 KB)
=====
downloaded 987 KB

OK
* Installing renv 0.12.0 ... Done!
Successfully installed and loaded renv 0.12.0.
* Project '/cloud/project/class-code' loaded. [renv 0.12.0]
* The project may be out of sync -- use `renv::status()` for more details.
> renv::status()
```

**Environment Tab:**

Global Environment

Environment is empty

**Files Tab:**

| Name             | Size    | Modified              |
|------------------|---------|-----------------------|
| ..               |         |                       |
| .Rprofile        | 26 B    | Oct 20, 2020, 1:48 PM |
| class-code.Rproj | 205 B   | Oct 20, 2020, 1:52 PM |
| code-here.R      | 101 B   | Oct 20, 2020, 1:48 PM |
| renv             |         |                       |
| renv.lock        | 17.6 KB | Oct 20, 2020, 1:48 PM |

# Install packages (then press 'y')

The screenshot shows the RStudio Cloud interface with the following components:

- Console Tab:** Displays the command `library("renv")` followed by a list of installed packages and their versions. The list includes: hms [0.5.3], htmltools [0.5.0], httr [1.4.2], isolbind [0.2.2], jsonlite [1.7.1], knitr [1.30], labeling [0.4.2], lattice [0.20-41], lifecycle [0.2.0], lubridate [1.7.9], magrittr [1.5], markdown [1.1], mvcv [1.8-33], mime [0.9], modelr [0.1.8], musell [0.5.0], nime [3.1-149], openssl [1.4.3], pillar [1.4.6], pkgbuild [1.1.0], pkgconfig [2.0.3], pkglload [1.1.0], praise [1.0.0], prettyunits [1.1.1], processx [3.4.4], progress [1.2.2], ps [1.4.0], purrr [0.3.4], readr [1.4.0], readxl [1.3.1], rematch [1.0.1], remotes [2.2.0], reprex [0.3.0], rmarkdown [2.4], rprojroot [1.3-2], rstudioapi [0.11], rvest [0.3.6], scales [1.1.1], selectr [0.4-2], stringi [1.5.3], stringr [1.4.0], sys [3.4], testthat [2.3.2], tibble [3.0.4], tidyR [1.1.2], tidyselect [1.1.0], tinytex [0.26], utf8 [1.1.4], vctrs [0.3.4], viridisLite [0.3.0], whisker [0.4], withr [2.1.0], xfun [0.18], xml2 [1.3.2], yaml [2.2.1].

Use `renv::snapshot()` to remove them from the lockfile.  
The following package(s) are recorded in the lockfile but not installed:  
nycflights13 [1.0.1]  
pacman [0.5.1]  
rlang [0.4.8]  
tidyverse [1.3.0]

Use `renv::restore()` to install these packages.

> `renv::restore()`
- Environment Tab:** Shows the message "Environment is empty".
- File Browser:** Shows the project structure in the "Cloud > project > class-code" directory. It includes files: .Rprofile (26 B, Oct 20, 2020, 1:48 PM), class-code.Rproj (205 B, Oct 20, 2020, 1:52 PM), code-here.R (101 B, Oct 20, 2020, 1:48 PM), renv (17.6 KB, Oct 20, 2020, 1:48 PM), and renv.lock.

# Quick intro to R

---

# Arithmetic operations

R can do all the standard arithmetic operations

```
1+2 ## add
```

```
## [1] 3
```

```
6-7 ## subtract
```

```
## [1] -1
```

```
5/2 ## divide
```

```
## [1] 2.5
```

# Logical operations

You also have logical operations

```
1 > 2
```

```
## [1] FALSE
```

```
(1 > 2) | (1 > 0.5) # / is the or operator
```

```
## [1] TRUE
```

```
(1 > 2) & (1 > 0.5) # & is the and operator
```

```
## [1] FALSE
```

# Logical operations

We can negate expressions with: !

This is helpful for filtering data

```
is.na(1:10)
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
!is.na(1:10)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

# Logical operators

For value matching we use: `%in%`

To see whether an object is contained within (i.e. matches one of) a list of items, use `%in%`.

```
4 %in% 1:10
```

```
## [1] TRUE
```

```
4 %in% 5:10
```

```
## [1] FALSE
```

This is kind of like an `any` command in other languages

# Logical operators

Turns out there's no `%notin%`, but we can make one using Negate

# Logical operators

Turns out there's no `%notin%`, but we can make one using `Negate`

```
`%notin%` ← Negate(`%in%`) ## The backticks (`) help to specify functions.  
4 %notin% 5:10
```

```
## [1] TRUE
```

# Logical operators

To evaluate whether two expressions are equal, we need to use **two** equal signs

```
1 = 1 ## This doesn't work
```

```
## Error in 1 = 1: invalid (do_set) left-hand side to assignment
```

```
1 == 1 ## This does.
```

```
## [1] TRUE
```

```
1 != 2 ## Note the single equal sign when combined with a negation.
```

```
## [1] TRUE
```

# Assignment

In R, we can use either `=` or `←` to handle assignment.<sup>1</sup>

<sup>1</sup> The `←` is really a `<` followed by a `-`. It just looks like an arrow because of the font on the slides.

# Assignment

In R, we can use either `=` or `←` to handle assignment.<sup>1</sup>

`←` is normally read aloud as "gets"

You can think of it as a (left-facing) arrow saying **assign in this direction**

<sup>1</sup> The `←` is really a `<` followed by a `-`. It just looks like an arrow because of the font on the slides.

# Assignment

```
a ← 10 + 5  
a
```

```
## [1] 15
```

# Assignment

```
a ← 10 + 5  
a
```

```
## [1] 15
```

An arrow can point in the other direction too (i.e. →), this is used less frequently though

```
10 + 5 → a  
a
```

```
## [1] 15
```

# Assignment

You can also use `=` for assignment, but the object must be on the left

```
b = 10 + 10  
b
```

```
## [1] 20
```

# Which assignment operator to use?

Most R folks prefer `←` for assignment

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

# Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions, we don't want to mix them up

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

# Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions, we don't want to mix them up

It doesn't really matter though, other languages use `=` for both

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

# Which assignment operator to use?

Most R folks prefer `←` for assignment

`=` has a specific role for evaluation *within* functions, we don't want to mix them up

It doesn't really matter though, other languages use `=` for both

Use whatever you prefer, just be consistent

In RStudio you can assign `←` to a hotkey to make using it as easy as `=`

# Help

If you are struggling with a (named) function or object in R, simply type ?

command here

```
?Negate
```

# Help

Also try `vignette()` for a more detailed introduction to many packages

```
# Try this:  
vignette("dplyr")
```

# Help

Also try `vignette()` for a more detailed introduction to many packages

```
# Try this:  
vignette("dplyr")
```

Vignettes are a very easy way to learn how and when to use a package

# Object-oriented programming

In R:

"Everything is an object and everything has a name."

# What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

# What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

A lot of these are probably familiar if you have coding experience

# What are objects?

We won't go into OOP details but here are some objects that we'll be working with regularly:

- vectors
- matrices
- data frames
- lists
- functions
- etc.

A lot of these are probably familiar if you have coding experience

But there are always language-specific features/subtleties

# Global environment

```
## Create a small data frame called "df".  
df <- data.frame(x = 1:2, y = 3:4)  
df
```

```
##   x y  
## 1 1 3  
## 2 2 4
```

# Global environment

```
## Create a small data frame called "df".  
df <- data.frame(x = 1:2, y = 3:4)  
df
```

```
##   x y  
## 1 1 3  
## 2 2 4
```

Now, let's try to run a regression<sup>1</sup> on these "x" and "y" variables:

<sup>1</sup> Yes, this is a dumb regression with perfectly co-linear variables. Just go with it.

# Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

# Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

```
## Error in eval(predvars, data, env): object 'y' not found
```

# Global environment

```
lm(y ~ x) ## The "lm" stands for linear model(s)
```

```
## Error in eval(predvars, data, env): object 'y' not found
```

Error?

# Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

# Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

R can't find the variables that we've supplied in our **Global Environment**

# Global environment

The error message is

```
## Error in eval(predvars, data, env): object 'y' not found
```

R can't find the variables that we've supplied in our **Global Environment**

Can you find x or y in the RStudio panel?

# Global environment

We have to tell R `x` and `y` are a part of the object `df`

# Global environment

We have to tell R `x` and `y` are a part of the object `df`

How?

# Global environment

We have to tell R `x` and `y` are a part of the object `df`

How?

There are a various ways to solve this problem. One is to simply specify the datasource:

```
lm(y ~ x, data = df) ## Works when we add "data = df"!
```

```
##  
## Call:  
## lm(formula = y ~ x, data = df)  
##  
## Coefficients:  
## (Intercept)          x  
##             2             1
```

# Global environment: why it matters

This matters largely for Stata users

# Global environment: why it matters

This matters largely for Stata users

In Stata, the workspace is basically just a single data frame  $\Rightarrow$  all variables are in the global environment

# Global environment: why it matters

This matters largely for Stata users

In Stata, the workspace is basically just a single data frame  $\Rightarrow$  all variables are in the global environment

Big problem with this is you can't have multiple data frames / datasets in memory

# Working with multiple objects

We can create a second data frame in memory!

```
df2 ← data.frame(x = rnorm(10), y = runif(10))  
df
```

```
##   x  y  
## 1 1  3  
## 2 2  4
```

```
df2
```

```
##           x          y  
## 1 -0.17584104 0.45717268  
## 2 -0.29812671 0.44376732  
## 3  1.31334913 0.06100947  
## 4  0.43630311 0.99799165  
## 5  0.51176927 0.45238594  
## 6 -0.77001834 0.05871020  
## 7 -0.22662226 0.32756661  
## 8  0.14535256 0.53055555  
## 9  0.62500000 0.12345678  
## 10 -0.05000000 0.77861278
```

# Reserved words

R has a bunch of key/reserved words that serve specific functions

See [here](#) for a full list, including (but not limited to):

```
if  
else  
while # looping  
function  
for # looping  
TRUE  
FALSE  
NULL # null/undefined  
Inf # infinity  
NaN # Not a number  
NA # Not available / missing
```

# Semi-reserved words

There are other words that are sort of reserved, in that they have a particular meaning

The most important one is `c()` which binds and concatenates objects together

```
my_vector ← c(1, 2, 5)  
my_vector
```

```
## [1] 1 2 5
```

Other ones are `pi`, `e`, etc

# Namespace conflicts

Try loading up `dplyr` in RStudio

```
library(dplyr)
```

What warning gets reported?

# Namespace conflicts

Try loading up `dplyr` in RStudio

```
library(dplyr)
```

What warning gets reported?

The warning *masked from 'package:X'* is about a **namespace conflict**

# Namespace conflicts

Try loading up `dplyr` in RStudio

```
library(dplyr)
```

What warning gets reported?

The warning *masked from 'package:X'* is about a **namespace conflict**

`dplyr` and the `stats` package (which gets loaded automatically when you start R) have functions named `filter` and `lag`

# Namespace conflicts

Whenever a namespace conflict arises, the most recently loaded package will gain preference

# Namespace conflicts

Whenever a namespace conflict arises, the most recently loaded package will gain preference

The `filter()` function now refers specifically to the `dplyr` variant

# Namespace conflicts

Whenever a namespace conflict arises, the most recently loaded package will gain preference

The `filter()` function now refers specifically to the `dplyr` variant

But what if we want the `stats` variant?

# Namespace conflicts

Whenever a namespace conflict arises, the most recently loaded package will gain preference

The `filter()` function now refers specifically to the `dplyr` variant

But what if we want the `stats` variant?

1. Temporarily use `stats::filter()`
2. Permanently assign `filter <- stats::filter`

# Solving namespace conflicts

Use `package::function()`

# Solving namespace conflicts

Use `package::function()`

Explicitly call a conflicted function from a package using the `package::function()` syntax:

```
stats::filter(1:10, rep(1, 2))
```

```
## Time Series:  
## Start = 1  
## End = 10  
## Frequency = 1  
## [1] 3 5 7 9 11 13 15 17 19 NA
```

# Solving namespace conflicts

We can also use `::` for more than just conflicted cases.

# Solving namespace conflicts

We can also use `::` for more than just conflicted cases.

It can clarify where a function or dataset comes from and make the code clearer:

```
dplyr::starwars ## Print the starwars data frame from the dplyr package  
scales::comma(c(1000, 1000000)) ## Use the comma function, which comes from the scales package
```

# Solving namespace conflicts

We can also use `::` for more than just conflicted cases.

It can clarify where a function or dataset comes from and make the code clearer:

```
dplyr::starwars ## Print the starwars data frame from the dplyr package  
scales::comma(c(1000, 1000000)) ## Use the comma function, which comes from the scales package
```

The `::` syntax also means that we can call functions without loading package first. E.g. As long as `dplyr` is installed on our system, then

```
dplyr::filter(iris, Species="virginica") will work.
```

# Solving namespace conflicts

Assign `function ← package::function`

# Solving namespace conflicts

```
Assign function <- package::function
```

You can permanently assign a conflicted function name to a particular package

# Solving namespace conflicts

```
Assign function <- package::function
```

You can permanently assign a conflicted function name to a particular package

This will hold for the remainder of your current R session, or until you change it back:

```
filter <- stats::filter ## Note the lack of parentheses.  
filter <- dplyr::filter ## Change it back again.
```

# Indexing

How do we index in R?

# Indexing

How do we index in R?

We've already seen an example of indexing in the form of R console output:

```
1+2
```

```
## [1] 3
```

The `[1]` above denotes the first (and, in this case, only) element of our output.<sup>1</sup>

# Indexing

How do we index in R?

We've already seen an example of indexing in the form of R console output:

```
1+2
```

```
## [1] 3
```

The `[1]` above denotes the first (and, in this case, only) element of our output.<sup>1</sup>

In this case, a vector of length one equal to the value "3"

# Indexing

Try the following in your console to see a more explicit example of indexed output:

```
rnorm(n = 100, mean = 0, sd = 1)
```

```
## [1]  0.720759329 -0.771380239  0.135003726 -1.275761215  0.893886090
## [6]  1.010140224 -1.785074795 -0.555924281  0.055263658  0.143093139
## [11] -0.276844843  0.151265833 -0.043763598 -0.576827912 -0.882678525
## [16] -0.913715522 -1.473331931  0.431167148  0.029567315  1.555507659
## [21]  1.382358798  0.111594350 -0.307941356 -0.050796507 -0.993041477
## [26] -1.175924112  1.650293519 -0.685798310 -1.018736464 -1.674876559
## [31] -1.036163414 -0.268123085 -0.323685890  0.024490509  0.243881278
## [36]  0.292811357 -0.861221473  2.253767108  0.243692154 -1.539791975
## [41] -0.272002112 -0.602838760 -1.412769793 -0.716187468 -1.279967975
## [46] -1.556153282 -1.317827138 -0.238362577 -1.844196614 -0.048207147
## [51]  0.698784951  0.034263852  1.116720329 -0.613634413  1.171298465
## [56] -1.075164463  0.487253184  0.585579872 -0.650666412 -0.986340053
## [61] -0.217147967 -0.700709072 -1.453705029 -1.312851083 -0.556284773
## [66] -0.391905133  0.345162388 -0.257688522  0.785867429 -0.604893591
```

# Indexing: []

We can also use [] to index objects that we create in R.

```
a <- 11:20  
a
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
a[4] ## Get the 4th element of object "a"
```

```
## [1] 14
```

```
a[c(4, 6)] ## Get the 4th and 6th elements
```

```
## [1] 14 16
```

# Indexing: []

It also works on larger arrays (vectors, matrices, data frames, and lists). For example:

```
starwars[1, 1] ## Show the cell corresponding to the 1st row & 1st column of the data frame.
```

```
## # A tibble: 1 x 1
##   name
##   <chr>
## 1 Luke Skywalker
```

# Indexing: []

It also works on larger arrays (vectors, matrices, data frames, and lists). For example:

```
starwars[1, 1] ## Show the cell corresponding to the 1st row & 1st column of the data frame.
```

```
## # A tibble: 1 x 1
##   name
##   <chr>
## 1 Luke Skywalker
```

What does `starwars[1:3, 1]` give you?

## Indexing: []

We haven't covered them properly yet (patience), but **lists** are a more complex type of array object in R

## Indexing: []

We haven't covered them properly yet (patience), but **lists** are a more complex type of array object in R

They can contain a random assortment of objects that don't share the same characteristics

# Indexing: []

We haven't covered them properly yet (patience), but **lists** are a more complex type of array object in R

They can contain a random assortment of objects that don't share the same characteristics

- e.g. a list can contain a scalar, a string, and a data frame, or even another list

# Indexing: []

The relevance to indexing is that lists require two square brackets [[]] to index the parent list item and then the standard [] within that parent item:

```
my_list <- list(  
  a = "hello",  
  b = c(1,2,3),  
  c = data.frame(x = 1:5, y = 6:10)  
)  
my_list[[1]] ## Return the 1st list object
```

```
## [1] "hello"
```

```
my_list[[2]][3] ## Return the 3rd element of the 2nd list object
```

```
## [1] 3
```

# Indexing: \$

Lists provide a nice segue to our other indexing operator: \$

- Let's continue with the my\_list example from the previous slide.

```
my_list
```

```
## $a  
## [1] "hello"  
##  
## $b  
## [1] 1 2 3  
##  
## $c  
##   x   y  
## 1 1   6  
## 2 2   7  
## 3 3   8  
## 4 4   9  
## 5 5  10
```

# Indexing: \$

Lists provide a nice segue to our other indexing operator: \$.

- Let's continue with the `my_list` example

```
my_list
```

```
## $a
## [1] "hello"
##
## $b
## [1] 1 2 3
##
## $c
##   x  y
## 1 1  6
## 2 2  7
## 3 3  8
## 4 4  9
## 5 5 10
```

# Indexing: \$

We can call these objects directly by name using the dollar sign, e.g.

```
my_list$a ## Return list object "a"
```

```
## [1] "hello"
```

```
my_list$b[3] ## Return the 3rd element of list object "b"
```

```
## [1] 3
```

```
my_list$c$x ## Return column "x" of list object "c"
```

```
## [1] 1 2 3 4 5
```

# Indexing: \$

The `$` form of indexing also works for other object types

In some cases, you can also combine the two index options:

```
starwars$name[1] # first element of the name column of the starwars data frame
```

```
## [1] "Luke Skywalker"
```

# Indexing: \$

However, note some key differences between the output from this example and that of our previous `starwars[1, 1]` example:

```
starwars$name[1]
```

```
## [1] "Luke Skywalker"
```

```
starwars[1, 1]
```

```
## # A tibble: 1 × 1
##   name
##   <chr>
## 1 Luke Skywalker
```

# Removing objects

Use `rm()` to remove an object or objects from your working environment.

```
a <- "hello"  
b <- "world"  
rm(a, b)
```

You can also use `rm(list = ls())` to remove all objects in your working environment (except packages), but this is **frowned upon**

# Removing objects

Use `rm()` to remove an object or objects from your working environment.

```
a <- "hello"  
b <- "world"  
rm(a, b)
```

You can also use `rm(list = ls())` to remove all objects in your working environment (except packages), but this is **frowned upon**

Just start a new R session instead

# The tidyverse

---

# What is "tidy" data?

## Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

# What is "tidy" data?

Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

# What is "tidy" data?

Resources:

- [Vignette](#) (from the `tidyr` package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Basically, tidy data is more likely to be [long \(i.e. narrow\)](#) than wide

# Checklist

Install tidyverse: `install.packages('tidyverse')`

Install nycflights13: `install.packages('nycflights13', repos = 'https://cran.rstudio.com')`

# Tidyverse vs. base R

Lots of debate over tidyverse vs base R

# Tidyverse vs. base R

Lots of debate over tidyverse vs base R

The answer is **obvious**: We should teach the tidyverse first

- Good documentation and support
- Consistent philosophy and syntax
- Nice front-end for big data tools
- For data cleaning, plotting, the tidyverse is elite

# Tidyverse vs. base R

Base R is still great

- Base R is extremely flexible and powerful
- The tidyverse can't do everything
- Using base R and the tidyverse together is often a good idea

# Tidyverse vs. base R

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent:

| tidyverse        | base                |
|------------------|---------------------|
| ?readr::read_csv | ?utils::read.csv    |
| ?dplyr::if_else  | ?base::ifelse       |
| ?tibble::tibble  | ?base :: data.frame |

Tidyverse functions typically have extra features on top of base R

# Tidyverse vs. base R

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent:

| tidyverse        | base                |
|------------------|---------------------|
| ?readr::read_csv | ?utils::read.csv    |
| ?dplyr::if_else  | ?base::ifelse       |
| ?tibble::tibble  | ?base :: data.frame |

Tidyverse functions typically have extra features on top of base R

There are always many ways to achieve a single goal in R

# Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

# Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We have actually loaded a number of packages: **ggplot2**, **tibble**, **dplyr**, etc

# Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We have actually loaded a number of packages: **ggplot2**, **tibble**, **dplyr**, etc

We can also see information about the package versions and some  
**namespace conflicts**

# Tidyverse packages

The tidyverse actually comes with a lot more packages than those that are just loaded automatically

```
tidyverse_packages()
```

```
## [1] "broom"        "cli"          "crayon"       "dbplyr"       "dplyr"  
## [6] "forcats"      "ggplot2"       "haven"        "hms"         "httr"  
## [11] "jsonlite"     "lubridate"    "magrittr"     "modelr"      "pillar"  
## [16] "purrr"        "readr"        "readxl"       "reprex"      "rlang"  
## [21] "rstudioapi"   "rvest"        "stringr"     "tibble"      "tidyverse"  
## [26] "xml2"         "tidyverse"
```

e.g. the **lubridate** package is for working with dates and the **rvest** package is for webscraping

# Tidyverse packages

The tidyverse actually comes with a lot more packages than those that are just loaded automatically

```
tidyverse_packages()
```

```
## [1] "broom"        "cli"          "crayon"       "dbplyr"       "dplyr"  
## [6] "forcats"      "ggplot2"       "haven"        "hms"         "httr"  
## [11] "jsonlite"     "lubridate"    "magrittr"     "modelr"      "pillar"  
## [16] "purrr"        "readr"        "readxl"       "reprex"      "rlang"  
## [21] "rstudioapi"   "rvest"        "stringr"     "tibble"      "tidyverse"  
## [26] "xml2"         "tidyverse"
```

e.g. the **lubridate** package is for working with dates and the **rvest** package is for webscraping

These packages have to be loaded separately

# Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

# Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

They are thus the ones that you will likely make the most use of

# Tidyverse packages

We're going to focus on two workhorse packages:

1. **dplyr**
2. **tidyr**

These are the packages for cleaning and wrangling data

They are thus the ones that you will likely make the most use of

Data cleaning and wrangling occupies an inordinate amount of time, no matter where you are in your research career

# Pipes: %>%

The pipe operator `%>%` lets us perform a sequence of operations in a very nice and tidy way

# Pipes: %>%

The pipe operator `%>%` lets us perform a sequence of operations in a very nice and tidy way

Suppose we wanted to figure out the average highway miles per gallon of Audi's in the `mpg` dataset:

```
mpg
```

```
## # A tibble: 234 x 11
##   manufacturer model     displ  year   cyl trans  drv   cty   hwy fl class
##   <chr>        <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi         a4      1.8  1999     4 auto(l... f       18    29 p     comp...
## 2 audi         a4      1.8  1999     4 manual... f      21    29 p     comp...
## 3 audi         a4      2.0  2008     4 manual... f      20    31 p     comp...
## 4 audi         a4      2.0  2008     4 auto(a... f      21    30 p     comp...
## 5 audi         a4      2.8  1999     6 auto(l... f      16    26 p     comp...
## 6 audi         a4      2.8  1999     6 manual... f      18    26 p     comp...
```

# Pipes: %>%

There's two ways you might do this without taking advantage of pipes:

# Pipes: %>%

There's two ways you might do this without taking advantage of pipes:

The first is to do it step-by-step, line-by-line which requires a lot of variable assignment

```
audis_mpg <- filter(mpg, manufacturer == "audi")
audis_mpg_grouped <- group_by(filter(mpg, manufacturer == "audi"), model)
summarise(audis_mpg_grouped, hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

# Pipes: %>%

Next you could do it all in one line which is hard to read

```
summarise(group_by(filter(mpg, manufacturer="audi")), model), hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

# Pipes: %>%

Or, you could use **pipes** %>%:

```
mpg %>% filter(manufacturer="audi") %>% group_by(model) %>% summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

# Pipes: %>%

Or, you could use **pipes** %>%:

```
mpg %>% filter(manufacturer="audi") %>% group_by(model) %>% summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>       <dbl>
## 1 a4          28.3
## 2 a4 quattro  25.8
## 3 a6 quattro  24
```

It performs the operations from left to right, exactly like you'd think of them:  
take this object (mpg), do this (filter), then do this (group by car model), then  
do this (take the mean of highway miles)

# Use vertical space

Pipes are even more readable if we write it over several lines:

```
mpg %>%
  filter(manufacturer=="audi") %>%
  group_by(model) %>%
  summarise(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 x 2
##   model      hwy_mean
##   <chr>        <dbl>
## 1 a4            28.3
## 2 a4 quattro    25.8
## 3 a6 quattro    24
```

Using vertical space costs nothing and makes for much more readable code

# dplyr

---

# Aside: dplyr 1.0.0 release

Please make sure that you are running at least **dplyr** 1.0.0 before continuing.

```
packageVersion("dplyr")
```

```
## [1] '1.0.2'
```

```
# install.packages('dplyr') ## install updated version if < 1.0.0
```

# The five key dplyr verbs

1. `filter`: Subset/filter rows based on their values
2. `arrange`: Reorder/arrange rows based on their values
3. `select`: Select columns/variables
4. `mutate`: Create new columns/variables
5. `summarise`: Collapse multiple rows into a single summary value, potentially by a grouping variable

# The five key dplyr verbs

1. `filter`: Subset/filter rows based on their values
2. `arrange`: Reorder/arrange rows based on their values
3. `select`: Select columns/variables
4. `mutate`: Create new columns/variables
5. `summarise`: Collapse multiple rows into a single summary value, potentially by a grouping variable

Let's practice these commands together using the `starwars` data frame that comes pre-packaged with `dplyr`

# Starwars

Here's the `starwars` dataset, it has 87 observations of 14 variables

```
starwars
```

```
## # A tibble: 87 x 14
##   name    height  mass hair_color skin_color eye_color birth_year sex gender
##   <chr>     <int> <dbl> <chr>       <chr>       <chr>          <dbl> <chr> <chr>
## 1 Luke...     172     77 blond      fair        blue            19 male   mascul...
## 2 C-3PO       167     75 <NA>       gold        yellow         112 none   mascul...
## 3 R2-D2        96     32 <NA>       white, bl... red           33 none   mascul...
## 4 Dart...      202    136 none      white        yellow         41.9 male   mascul...
## 5 Leia...      150     49 brown      light       brown           19 fema... femin...
## 6 Owen...      178    120 brown, gr... light       blue            52 male   mascul...
## 7 Beru...      165     75 brown      light       blue            47 fema... femin...
## 8 R5-D4        97     32 <NA>       white, red red            NA none   mascul...
## 9 Bigg...      183     84 black      light       brown           24 male   mascul...
## 10 Obi-...      182     77 auburn, w... fair        blue-gray        57 male   mascul...
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

# 1) dplyr::filter

Here we are subsetting the observations of humans that are at least 190cm

```
starwars %>%  
  filter(  
    species = "Human",  
    height >= 190  
)
```

```
## # A tibble: 4 x 14  
##   name  height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>   <int> <dbl> <chr>       <chr>       <chr>          <dbl> <chr> <chr>  
## 1 Dart...     202    136 none      white       yellow        41.9 male   masculin...  
## 2 Qui-...     193     89 brown     fair        blue         92   male   masculin...  
## 3 Dooku      193     80 white     fair        brown        102  male   masculin...  
## 4 Bail...     191     NA black     tan         brown        67   male   masculin...  
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

# 1) dplyr::filter

You can filter using regular expressions with grep-type commands or the `stringr` package

```
starwars %>%  
  filter(stringr::str_detect(name, "Skywalker"))
```

```
## # A tibble: 3 x 14  
##   name    height   mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>    <int>  <dbl> <chr>      <chr>      <chr>        <dbl> <chr> <chr>  
## 1 Luke...     172     77 blond     fair       blue         19 male   masculin...  
## 2 Anak...     188     84 blond     fair       blue        41.9 male   masculin...  
## 3 Shmi...     163     NA black     fair       brown        72 female feminin...
```

## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## # vehicles <list>, starships <list>

This subsets the observations for individuals whose names contain "Skywalker"

# 1) dplyr::filter

A very common `filter` use case is identifying/removing missing data cases:

```
starwars %>%  
  filter(is.na(height))
```

```
## # A tibble: 6 x 14  
##   name  height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>   <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>  
## 1 Arve...     NA     NA brown      fair        brown         NA male  masculin...  
## 2 Finn       NA     NA black      dark        dark          NA male  masculin...  
## 3 Rey        NA     NA brown      light       hazel         NA femal... feminin...  
## 4 Poe ...     NA     NA brown      light       brown         NA male  masculin...  
## 5 BB8        NA     NA none       none       black         NA none  masculin...  
## 6 Capt...     NA     NA unknown    unknown    unknown        NA <NA>  <NA>  
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

# 1) dplyr::filter

To remove missing observations, use negation:

```
starwars %>%  
  filter(!is.na(height))
```

```
## # A tibble: 81 x 14  
##   name    height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>       <chr>       <chr>        <dbl> <chr> <chr>  
## 1 Luke...     172     77 blond      fair        blue          19 male  masculin...  
## 2 C-3PO       167     75 <NA>       gold        yellow        112 none  masculin...  
## 3 R2-D2        96     32 <NA>       white, bl... red           33 none  masculin...  
## 4 Dart...      202    136 none       white        yellow        41.9 male  masculin...  
## 5 Leia...      150     49 brown      light       brown          19 female feminin...  
## 6 Owen...      178    120 brown, gr... light       blue           52 male  masculin...  
## 7 Beru...      165     75 brown      light       blue           47 female feminin...  
## 8 R5-D4        97     32 <NA>       white, red red            NA none  masculin...  
## 9 Bigg...      183     84 black      light       brown          24 male  masculin...  
## 10 Obi-...      182     77 auburn, w... fair        blue-gray        57 male  masculin...  
## # ... with 71 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

## 2) dplyr::arrange

`arrange` sorts the data frame based on the variables you supply:

```
starwars %>%  
  arrange(birth_year)
```

```
## # A tibble: 87 x 14  
##   name   height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>    <int> <dbl> <chr>       <chr>       <chr>        <dbl> <chr> <chr>  
## 1 Wick...     88    20  brown      brown       brown          8   male  masculin...  
## 2 IG-88      200   140 none       metal       red            15  none  masculin...  
## 3 Luke...     172    77 blond      fair        blue           19  male  masculin...  
## 4 Leia...     150    49 brown      light       brown          19  femal... feminin...  
## 5 Wedg...     170    77 brown      fair        hazel          21  male  masculin...  
## 6 Plo ...     188    80 none       orange     black           22  male  masculin...  
## 7 Bigg...     183    84 black      light       brown          24  male  masculin...  
## 8 Han ...     180    80 brown      fair        brown          29  male  masculin...  
## 9 Land...     177    79 black      dark        brown          31  male  masculin...  
## 10 Boba...    183   78.2 black      fair        brown         31.5 male  masculin...  
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

## 2) dplyr::arrange

We can also arrange items in descending order using `arrange(desc())`

```
starwars %>%  
  arrange(desc(birth_year))
```

```
## # A tibble: 87 x 14  
##   name    height  mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>     <int> <dbl> <chr>       <chr>       <chr>        <dbl> <chr> <chr>  
## 1 Yoda      66     17  white       green       brown          896 male   masculin...  
## 2 Jabb...    175    1358 <NA>       green-tan... orange          600 herm... masculin...  
## 3 Chew...    228    112  brown       unknown      blue           200 male   masculin...  
## 4 C-3PO     167     75 <NA>       gold        yellow         112 none   masculin...  
## 5 Dooku     193     80  white       fair        brown          102 male   masculin...  
## 6 Qui-...    193     89  brown       fair        blue            92 male   masculin...  
## 7 Ki-A...    198     82  white       pale        yellow         92 male   masculin...  
## 8 Fini...    170     NA  blond       fair        blue           91 male   masculin...  
## 9 Palp...    170     75  grey        pale        yellow         82 male   masculin...  
## 10 Clie...   183     NA  brown       fair        blue           82 male   masculin...  
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

### 3) dplyr::select

Use commas to select multiple columns out of a data frame, deselect a column with "-", select across multiple columns with "first:last":

```
starwars %>%  
  select(name:skin_color, species, -height)
```

```
## # A tibble: 87 x 5  
##   name           mass hair_color   skin_color species  
##   <chr>        <dbl> <chr>       <chr>     <chr>  
## 1 Luke Skywalker    77  blond      fair       Human  
## 2 C-3PO              75  <NA>       gold      Droid  
## 3 R2-D2              32  <NA>       white, blue Droid  
## 4 Darth Vader       136  none       white      Human  
## 5 Leia Organa        49  brown      light      Human  
## 6 Owen Lars          120 brown, grey light      Human  
## 7 Beru Whitesun lars  75  brown      light      Human  
## 8 R5-D4              32  <NA>       white, red Droid  
## 9 Biggs Darklighter   84  black      light      Human  
## 10 Obi-Wan Kenobi     77  auburn, white fair      Human
```

### 3) dplyr::select

You can also rename your selected variables in place

```
starwars %>%  
  select(alias = name, crib = homeworld)
```

```
## # A tibble: 87 x 2  
##   alias      crib  
##   <chr>      <chr>  
## 1 Luke Skywalker Tatooine  
## 2 C-3PO        Tatooine  
## 3 R2-D2        Naboo  
## 4 Darth Vader Tatooine  
## 5 Leia Organa Alderaan  
## 6 Owen Lars   Tatooine  
## 7 Beru Whitesun lars Tatooine  
## 8 R5-D4        Tatooine  
## 9 Biggs Darklighter Tatooine  
## 10 Obi-Wan Kenobi Stewjon  
## # ... with 77 more rows
```

### 3) dplyr::select

If you just want to rename columns without subsetting them, you can use

rename:

```
starwars %>%  
  rename(alias = name, crib = homeworld)
```

```
## # A tibble: 87 x 14  
##   alias height mass hair_color skin_color eye_color birth_year sex gender  
##   <chr>    <int> <dbl> <chr>       <chr>       <chr>        <dbl> <chr> <chr>  
## 1 Luke...     172     77 blond      fair        blue          19 male  masculin...  
## 2 C-3PO       167     75 <NA>       gold        yellow        112 none  masculin...  
## 3 R2-D2        96     32 <NA>       white, bl... red           33 none  masculin...  
## 4 Dart...      202    136 none       white        yellow        41.9 male  masculin...  
## 5 Leia...      150     49 brown      light       brown          19 female feminin...  
## 6 Owen...      178    120 brown, gr... light       blue          52 male  masculin...  
## 7 Beru...      165     75 brown      light       blue          47 female feminin...  
## 8 R5-D4        97     32 <NA>       white, red red           NA none  masculin...  
## 9 Bigg...      183     84 black      light       brown          24 male  masculin...  
## 10 Obi-       182     77 auburn, w... fair        blue-gray       57 male  masculin...
```

### 3) dplyr::select cont.

The `select(contains(PATTERN))` option provides a nice shortcut in relevant cases.

```
starwars %>%  
  select(name, contains("color"))  
  
## # A tibble: 87 x 4  
##   name           hair_color   skin_color eye_color  
##   <chr>          <chr>       <chr>      <chr>  
## 1 Luke Skywalker    blond      fair        blue  
## 2 C-3PO              <NA>       gold        yellow  
## 3 R2-D2              <NA>       white, blue red  
## 4 Darth Vader       none       white       yellow  
## 5 Leia Organa        brown      light       brown  
## 6 Owen Lars          brown, grey light       blue  
## 7 Beru Whitesun lars brown      light       blue  
## 8 R5-D4              <NA>       white, red red  
## 9 Biggs Darklighter  black      light       brown  
## 10 Obi-Wan Kenobi    auburn     white, fair blue-gray
```

### 3) dplyr::select

The `select( ... , everything())` option is another useful shortcut if you only want to bring some variable(s) to the "front" of a data frame

```
starwars %>%  
  select(species, homeworld, everything()) %>%  
  head(5)  
  
## # A tibble: 5 x 14  
##   species homeworld name  height  mass hair_color skin_color eye_color  
##   <chr>     <chr>    <chr>   <int>   <dbl>   <chr>       <chr>       <chr>  
## 1 Human     Tatooine  Luke...     172     77  blond        fair        blue  
## 2 Droid      Tatooine  C-3PO      167     75 <NA>        gold        yellow  
## 3 Droid      Naboo     R2-D2      96      32 <NA>      white, bl... red  
## 4 Human     Tatooine  Dart...     202    136  none        white        yellow  
## 5 Human     Alderaan  Leia...     150     49  brown       light        brown  
## # ... with 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

### 3) dplyr::select

You can also use `relocate` to do the same thing

```
starwars %>%  
  relocate(species, homeworld) %>%  
  head(5)
```

```
## # A tibble: 5 x 14  
##   species homeworld name   height  mass hair_color skin_color eye_color  
##   <chr>     <chr>    <chr>   <int>   <dbl>   <chr>       <chr>       <chr>  
## 1 Human     Tatooine  Luke...     172     77  blond        fair        blue  
## 2 Droid      Tatooine  C-3PO      167     75 <NA>        gold        yellow  
## 3 Droid      Naboo     R2-D2      96      32 <NA>       white, bl... red  
## 4 Human     Tatooine  Dart...     202    136  none        white        yellow  
## 5 Human     Alderaan  Leia...     150     49  brown       light        brown  
## # ... with 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>,  
## #   films <list>, vehicles <list>, starships <list>
```

## 4) dplyr::mutate

You can create new columns from scratch as transformations of existing columns:

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(dog_years = birth_year * 7) %>%  
  mutate(comment = paste0(name, " is ", dog_years, " in dog years.))
```

```
## # A tibble: 87 x 4  
##   name           birth_year  dog_years comment  
##   <chr>          <dbl>      <dbl> <chr>  
## 1 Luke Skywalker     19        133  Luke Skywalker is 133 in dog years.  
## 2 C-3PO              112       784  C-3PO is 784 in dog years.  
## 3 R2-D2              33        231  R2-D2 is 231 in dog years.  
## 4 Darth Vader        41.9      293. Darth Vader is 293.3 in dog years.  
## 5 Leia Organa         19        133  Leia Organa is 133 in dog years.  
## 6 Owen Lars            52       364  Owen Lars is 364 in dog years.  
## 7 Beru Whitesun lars    47       329  Beru Whitesun lars is 329 in dog yea...  
## 8 R5-D4              NA        NA   R5-D4 is NA in dog years
```

## 4) dplyr::mutate

Note: `mutate` creates variables in order, so you can chain multiple mutates in a single call

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(  
    dog_years = birth_year * 7, ## Separate with a comma  
    comment = paste0(name, " is ", dog_years, " in dog years.")  
  )
```

```
## # A tibble: 87 x 4  
##   name      birth_year  dog_years comment  
##   <chr>        <dbl>     <dbl> <chr>  
## 1 Luke Skywalker      19       133  Luke Skywalker is 133 in dog years.  
## 2 C-3PO                 112      784  C-3PO is 784 in dog years.  
## 3 R2-D2                  33      231  R2-D2 is 231 in dog years.  
## 4 Darth Vader             41.9    293. Darth Vader is 293.3 in dog years.  
## 5 Leia Organa              19       133  Leia Organa is 133 in dog years.  
## 6 Owen Lars                52      364  Owen Lars is 364 in dog years.
```

## 4) dplyr::mutate

Boolean, logical and conditional operators all work well with `mutate` too:

```
starwars %>%
  select(name, height) %>%
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) %>%
  mutate(tall1 = height > 180) %>% # TRUE or FALSE
  mutate(tall2 = ifelse(height > 180, "Tall", "Short")) ## Same effect, but can choose labels
```

```
## # A tibble: 2 x 4
##   name           height tall1 tall2
##   <chr>          <int> <lgl> <chr>
## 1 Luke Skywalker     172 FALSE Short
## 2 Anakin Skywalker    188 TRUE  Tall
```

## 4) dplyr::mutate

Lastly, combining `mutate` with `across` allows you to easily work on a subset of variables:

```
starwars %>%  
  select(name:eye_color) %>%  
  mutate(across(where(is.character), toupper)) %>% # Take all character variables, uppercase them  
  head(5)
```

```
## # A tibble: 5 x 6  
##   name           height  mass hair_color skin_color eye_color  
##   <chr>        <int> <dbl> <chr>      <chr>      <chr>  
## 1 LUKE SKYWALKER     172    77 BLOND      FAIR       BLUE  
## 2 C-3PO              167    75 <NA>       GOLD       YELLOW  
## 3 R2-D2               96    32 <NA>      WHITE, BLUE  RED  
## 4 DARTH VADER       202   136 NONE       WHITE      YELLOW  
## 5 LEIA ORGANA        150    49 BROWN     LIGHT      BROWN
```

# 5) dplyr::summarise

Summarising useful in combination with the `group_by` command

```
starwars %>%  
  group_by(species, gender) %>% # for each species-gender combo  
  summarise(mean_height = mean(height, na.rm = TRUE)) # calculate the mean height
```

```
## # A tibble: 42 x 3  
## # Groups:   species [38]  
##   species   gender   mean_height  
##   <chr>     <chr>       <dbl>  
## 1 Aleena    masculine      79  
## 2 Besalisk  masculine     198  
## 3 Cerean    masculine     198  
## 4 Chagrian  masculine     196  
## 5 Clawdite  feminine      168  
## 6 Droid     feminine       96  
## 7 Droid     masculine     140  
## 8 Dug       masculine     112  
## 9 Ewok      masculine      88  
## 10 Geonosian masculine     182
```

## 5) dplyr::summarise

Note that including "na.rm = TRUE" is usually a good idea with summarise functions, it keeps NAs from propagating to the end result

```
## Probably not what we want
starwars %>%
  summarise(mean_height = mean(height))
```

```
## # A tibble: 1 x 1
##   mean_height
##       <dbl>
## 1        NA
```

## 5) dplyr::summarise

We can also use `across` within `summarise`:

```
starwars %>%  
  group_by(species) %>% # for each species  
  summarise(across(where(is.numeric), mean, na.rm = T)) %>% # take the mean of all numeric varian  
  head(5)
```

```
## # A tibble: 5 x 4  
##   species    height   mass birth_year  
##   <chr>      <dbl>   <dbl>     <dbl>  
## 1 Aleena       79     15       NaN  
## 2 Besalisk     198    102       NaN  
## 3 Cerean       198     82       92  
## 4 Chagrian     196     NaN       NaN  
## 5 Clawdite     168     55       NaN
```

# Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping

- Particularly useful with the `summarise` and `mutate` commands

# Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping

- Particularly useful with the `summarise` and `mutate` commands

`slice`: Subset rows by position rather than filtering by values

- E.g. `starwars %>% slice(c(1, 5))`

# Other dplyr goodies

`pull`: Extract a column from as a data frame as a vector or scalar

- E.g. `starwars %>% filter(gender=="female") %>% pull(height)`

# Other dplyr goodies

`pull`: Extract a column from as a data frame as a vector or scalar

- E.g. `starwars %>% filter(gender=="female") %>% pull(height)`

`count` and `distinct`: Number and isolate unique observations

- E.g. `starwars %>% count(species)`, or `starwars %>% distinct(species)`
- You could also use a combination of `mutate`, `group_by`, and `n()`, e.g.  
`starwars %>% group_by(species) %>% mutate(num = n())`.

# Other dplyr goodies

There are also a whole class of **window functions** for getting leads and lags, percentiles, cumulative sums, etc.

- See `vignette("window-functions")`.

## dplyr::xxxx\_join

The last set of commands we need are the `join` commands

## dplyr::xxxx\_join

The last set of commands we need are the `join` commands

These are the same as `merge` in stata but with a bit more functionality

# dplyr::xxxx\_join

We merge data with **join operations**:

- inner\_join(df1, df2)
- left\_join(df1, df2)
- right\_join(df1, df2)
- full\_join(df1, df2)
- semi\_join(df1, df2)
- anti\_join(df1, df2)

(You can visualize the operations [here](#))

# dplyr::xxxx\_join

Lets use the data that comes with the the [nycflights13](#) package.

```
library(nycflights13)
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>     <int>           <int>     <dbl>     <int>           <int>
## 1 2013     1     1      517             515        2     830            819
## 2 2013     1     1      533             529        4     850            830
## 3 2013     1     1      542             540        2     923            850
## 4 2013     1     1      544             545       -1    1004           1022
## 5 2013     1     1      554             600       -6    812            837
## 6 2013     1     1      554             558       -4    740            728
## 7 2013     1     1      555             600       -5    913            854
## 8 2013     1     1      557             600       -3    709            723
## 9 2013     1     1      557             600       -3    838            846
## 10 2013    1     1      558             600       -2    753            745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier_schd <fct>, flight_schd <dbl>, tailnum_schd <chr>, origin_schd <chr>,
## #   dest_schd <chr>, arr_time_schd <dbl>, dep_time_schd <dbl>, sched_dep_time_schd <dbl>,
## #   arr_delay_schd <dbl>, dep_delay_schd <dbl>, sched_arr_time_schd <dbl>
```

# dplyr::xxxx\_join

planes

```
## # A tibble: 3,322 x 9
##   tailnum year type      manufacturer    model engines seats speed engine
##   <chr>   <int> <chr>      <chr>        <chr>   <int> <int> <int> <chr>
## 1 N10156  2004 Fixed wing m... EMBRAER       EMB-1...     2     55   NA Turbo-...
## 2 N102UW   1998 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 3 N103US   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 4 N104UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 5 N10575   2002 Fixed wing m... EMBRAER       EMB-1...     2     55   NA Turbo-...
## 6 N105UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 7 N107US   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 8 N108UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 9 N109UW   1999 Fixed wing m... AIRBUS INDUST... A320-...     2    182   NA Turbo-...
## 10 N110UW  1999 Fixed wing m... AIRBUS INDUST... A320-...    2    182   NA Turbo-...
## # ... with 3,312 more rows
```

# Joining operations

Let's perform a left join on the flights and planes datasets

- Note: I'm going to subset columns after the join, but only to keep text on the slide

# Joining operations

Let's perform a left join on the flights and planes datasets

- Note: I'm going subset columns after the join, but only to keep text on the slide

```
left_join(flights, planes) %>%  
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
```

```
## Joining, by = c("year", "tailnum")
```

```
## # A tibble: 336,776 x 10  
##   year month   day dep_time arr_time carrier flight tailnum type model  
##   <int> <int> <int>    <int>    <int> <chr>    <int> <chr>    <chr> <chr>  
## 1 2013     1     1      517      830  UA        1545 N14228 <NA> <NA>  
## 2 2013     1     1      533      850  UA        1714 N24211 <NA> <NA>  
## 3 2013     1     1      542      923  AA        1141 N619AA <NA> <NA>  
## 4 2013     1     1      544     1004  B6        725  N804JB <NA> <NA>  
## 5 2013     1     1      554      812  DL        461  N668DN <NA> <NA>
```

# Joining operations

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name), and told us what it chose

```
## Joining, by = c("year", "tailnum")
```

There's an obvious problem here: the variable `year` does not have a consistent meaning across our joining datasets

# Joining operations

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name), and told us what it chose

```
## Joining, by = c("year", "tailnum")
```

There's an obvious problem here: the variable `year` does not have a consistent meaning across our joining datasets

In one it refers to the *year of flight*, in the other it refers to *year of construction*

Luckily, there's an easy way to avoid this problem: try `?dplyr :: join`

# Joining operations

You just need to be more explicit in your join call by using the `by =` argument

```
left_join(  
  flights,  
  planes %>% rename(year_built = year), ## Not necessary w/ below line, but helpful  
  by = "tailnum" ## Be specific about the joining column  
 ) %>%  
 select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, year_built, type, model)  
 head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 x 11  
##   year month   day dep_time arr_time carrier flight tailnum year_built type  
##   <int> <int> <int>    <int>    <int> <chr>    <int> <chr>      <int> <chr>  
## 1  2013     1     1      517      830  UA        1545 N14228      1999 Fixe...  
## 2  2013     1     1      533      850  UA        1714 N24211      1998 Fixe...  
## 3  2013     1     1      542      923  AA        1141 N619AA      1990 Fixe...  
## # ... with 1 more variable: model <chr>
```

# Joining operations

Note what happens if we again specify the join column but don't rename the ambiguous year:

```
left_join(flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum") %>%
  select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
  head(3)

## # A tibble: 3 x 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type   model
##   <int>  <int> <int> <int>    <int>    <int> <chr>   <int> <chr>   <chr> <chr>
## 1  2013    1999     1     1      517      830  UA        1545 N14228  Fixe... 737-...
## 2  2013    1998     1     1      533      850  UA        1714 N24211  Fixe... 737-...
## 3  2013    1990     1     1      542      923  AA        1141 N619AA  Fixe... 757-...
```

# Joining operations

Note what happens if we again specify the join column but don't rename the ambiguous year:

```
left_join(flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum") %>%
  select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
  head(3)

## # A tibble: 3 x 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type   model
##   <int>  <int> <int> <int>    <int>    <int> <chr>   <int> <chr>   <chr> <chr>
## 1  2013    1999     1     1      517      830  UA        1545 N14228  Fixe... 737-...
## 2  2013    1998     1     1      533      850  UA        1714 N24211  Fixe... 737-...
## 3  2013    1990     1     1      542      923  AA        1141 N619AA  Fixe... 757-...
```

Make sure you know what "year.x" and "year.y" are

# tidyr

---

# Key `tidyr` verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt", "reshape long")
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast", "reshape wide")
3. `separate`: Split one column into multiple columns
4. `unite`: Combine multiple columns into one

# Key tidy verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt", "reshape long")
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast", "reshape wide")
3. `separate`: Split one column into multiple columns
4. `unite`: Combine multiple columns into one

Let's practice these verbs together in class

# 1) tidy::pivot\_longer

```
stocks ← data.frame(  
  time = as.Date('2009-01-01') + 0:1,  
  X = rnorm(2, 0, 1),  
  Y = rnorm(2, 0, 2),  
  Z = rnorm(2, 0, 4)  
)  
stocks
```

```
##          time        X        Y        Z  
## 1 2009-01-01 -1.146916 -4.1891925  0.7006416  
## 2 2009-01-02 -1.491875  0.9525605 -8.1531928
```

We have 4 variables, the date and the stocks

How do we get this in tidy form?

# 1) tidyverse::pivot\_longer

```
stocks %>% pivot_longer(-time, names_to = "stock", values_to = "price")
```

We need to pivot the stock name variables X, Y, Z longer

1. Choose non-time variables: -time
2. Decide what variable holds the names: names\_to = "stock"
3. Decide what variable holds the values: values\_to = "price"

# 1) tidyverse::pivot\_longer

```
stocks %>% pivot_longer(-time, names_to = "stock", values_to = "price")
```

```
## # A tibble: 6 x 3
##   time     stock   price
##   <date>    <chr>   <dbl>
## 1 2009-01-01 X      -1.15
## 2 2009-01-01 Y      -4.19
## 3 2009-01-01 Z       0.701
## 4 2009-01-02 X      -1.49
## 5 2009-01-02 Y       0.953
## 6 2009-01-02 Z      -8.15
```

# 1) tidyverse::pivot\_longer

Let's quickly save the "tidy" (i.e. long) stocks data frame for use on the next slide

```
tidy_stocks ← stocks %>%  
  pivot_longer(-time, names_to = "stock", values_to = "price")
```

## 2) tidyverse::pivot\_wider

```
tidy_stocks %>% pivot_wider(names_from = stock, values_from = price)
```

```
## # A tibble: 2 × 4
##   time           X       Y       Z
##   <date>     <dbl>   <dbl>   <dbl>
## 1 2009-01-01 -1.15  -4.19   0.701
## 2 2009-01-02 -1.49   0.953  -8.15
```

```
tidy_stocks %>% pivot_wider(names_from = time, values_from = price)
```

```
## # A tibble: 3 × 3
##   stock `2009-01-01` `2009-01-02`
##   <chr>     <dbl>      <dbl>
## 1 X         -1.15      -1.49
## 2 Y         -4.19      0.953
## 3 Z          0.701     -8.15
```

## 2) tidyverse::pivot\_wider

```
tidy_stocks %>% pivot_wider(names_from = stock, values_from = price)
```

```
## # A tibble: 2 x 4
##   time           X       Y       Z
##   <date>     <dbl>   <dbl>   <dbl>
## 1 2009-01-01 -1.15 -4.19  0.701
## 2 2009-01-02 -1.49  0.953 -8.15
```

```
tidy_stocks %>% pivot_wider(names_from = time, values_from = price)
```

```
## # A tibble: 3 x 3
##   stock `2009-01-01` `2009-01-02`
##   <chr>     <dbl>      <dbl>
## 1 X         -1.15      -1.49
## 2 Y         -4.19      0.953
## 3 Z          0.701     -8.15
```

Note that the second example has effectively transposed the data

### 3) tidyverse::separate

```
economists ← data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))
economists
```

```
##           name
## 1      Adam.Smith
## 2  Paul.Samuelson
## 3 Milton.Friedman
```

```
economists %>% separate(name, c("first_name", "last_name"))
```

```
##   first_name last_name
## 1      Adam     Smith
## 2      Paul Samuelson
## 3    Milton   Friedman
```

### 3) tidyverse::separate

```
conomists ← data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))
conomists
```

```
##           name
## 1      Adam.Smith
## 2  Paul.Samuelson
## 3 Milton.Friedman
```

```
conomists %>% separate(name, c("first_name", "last_name"))
```

```
##   first_name last_name
## 1      Adam     Smith
## 2      Paul Samuelson
## 3    Milton   Friedman
```

This command is pretty smart. But to avoid ambiguity, you can also specify the separation character with `separate( ... , sep=".")`

### 3) `tidyr::separate`

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data)

```
jobs ← data.frame(  
  name = c("Jack", "Jill"),  
  occupation = c("Homemaker", "Philosopher, Philanthropist, Troublemaker")  
)  
jobs
```

```
##   name                      occupation  
## 1 Jack                      Homemaker  
## 2 Jill Philosopher, Philanthropist, Troublemaker
```

### 3) `tidyr::separate`

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data)

```
## Now split out Jill's various occupations into different rows
jobs %>% separate_rows(occupation)
```

```
## # A tibble: 4 x 2
##   name  occupation
##   <chr> <chr>
## 1 Jack   Homemaker
## 2 Jill   Philosopher
## 3 Jill   Philanthropist
## 4 Jill   Troublemaker
```

## 4) tidyverse

```
gdp ← data.frame(  
  yr = rep(2016, times = 4),  
  mnth = rep(1, times = 4),  
  dy = 1:4,  
  gdp = rnorm(4, mean = 100, sd = 2)  
)  
gdp
```

```
##      yr mnth dy      gdp  
## 1 2016     1   1 99.74637  
## 2 2016     1   2 102.59236  
## 3 2016     1   3 100.50825  
## 4 2016     1   4  97.95164
```

## 4) tidyverse

```
## Combine "yr", "mnth", and "dy" into one "date" column  
gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-")
```

```
##          date      gdp  
## 1 2016-1-1  99.74637  
## 2 2016-1-2 102.59236  
## 3 2016-1-3 100.50825  
## 4 2016-1-4  97.95164
```

## 4) tidyverse::unite

Note that `unite` will automatically create a character variable:

```
gdp_u <- gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-") %>% as_tibble()  
gdp_u
```

```
## # A tibble: 4 x 2  
##   date      gdp  
##   <chr>     <dbl>  
## 1 2016-1-1  99.7  
## 2 2016-1-2  103.  
## 3 2016-1-3  101.  
## 4 2016-1-4  98.0
```

## 4) `tidyr::unite`

Note that `unite` will automatically create a character variable:

```
gdp_u ← gdp %>% unite(date, c("yr", "mnth", "dy"), sep = "-") %>% as_tibble()  
gdp_u
```

```
## # A tibble: 4 x 2  
##   date      gdp  
##   <chr>     <dbl>  
## 1 2016-1-1  99.7  
## 2 2016-1-2  103.  
## 3 2016-1-3  101.  
## 4 2016-1-4  98.0
```

If you want to convert it to something else (e.g. date or numeric) then you will need to modify it using `mutate`

## 4) tidyverse

```
library(lubridate)
gdp_u %>% mutate(date = ymd(date))
```

```
## # A tibble: 4 x 2
##   date      gdp
##   <date>    <dbl>
## 1 2016-01-01  99.7
## 2 2016-01-02 103.
## 3 2016-01-03 101.
## 4 2016-01-04  98.0
```

# Other tidyverse goodies

Use `crossing` to get the full combination of a group of variables

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 x 2
##   side  height
##   <chr> <chr>
## 1 left   bottom
## 2 left   top
## 3 right  bottom
## 4 right  top
```

# Other tidyverse goodies

Use `crossing` to get the full combination of a group of variables

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 x 2
##   side  height
##   <chr> <chr>
## 1 left   bottom
## 2 left   top
## 3 right  bottom
## 4 right  top
```

See `?expand` and `?complete` for more specialized functions that allow you to fill in (implicit) missing data or variable combinations in existing data frames

# Randomization and inference

---

# The Rubin causal model

**What is a causal effect?**

# The Rubin causal model

## What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

# The Rubin causal model

## What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

# The Rubin causal model

## What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

e.g. what is the difference in health between high and low levels of air pollution?

# The Rubin causal model

## What is a causal effect?

A starting point is that a causal effect is a **comparison between two potential outcomes**

What is the difference in some outcome in the presence vs the absence of a given treatment?

e.g. what is the difference in health between high and low levels of air pollution?

Let's begin formalizing this idea

# The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

# The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

The potential outcome is  $Y_i^1$  if unit  $i$  received some treatment, and  $Y_i^0$  if the unit did not

# The Rubin causal model: potential outcomes

Suppose we have a set of **observational units**

- These can be people, states, animals, air quality monitors, etc

Each unit has two **potential outcomes**, but only one is observed

The potential outcome is  $Y_i^1$  if unit  $i$  received some treatment, and  $Y_i^0$  if the unit did not

$Y_i^0$  corresponds to the control state of the world for  $i$

# The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

# The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

Observable outcomes are outcomes that actually show up in the data (i.e. the *factual outcome*)<sup>1</sup>

<sup>1</sup>The potential outcome that did not happen is called the *counterfactual outcome*.

# The Rubin causal model: observable outcomes

Note that these potential outcomes are not the same as observable outcomes

Observable outcomes are outcomes that actually show up in the data (i.e. the *factual outcome*)<sup>1</sup>

We can write the observable outcome  $Y_i$  as a simple equation:

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

where  $D_i = 1$  if  $i$  received treatment and 0 otherwise

<sup>1</sup>The potential outcome that did not happen is called the *counterfactual outcome*.

# The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

# The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

Rubin defines a treatment/causal effect  $\delta_i$  as the difference between the two potential outcomes:

$$\delta_i = Y_i^1 - Y_i^0$$

# The Rubin causal model: treatment effects

$$Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$$

This is the basis of the *Rubin causal model*

Rubin defines a treatment/causal effect  $\delta_i$  as the difference between the two potential outcomes:

$$\delta_i = Y_i^1 - Y_i^0$$

This leads to an obvious problem: we only observe one of these two states for each unit  $i$  but we need to know both to recover  $\delta_i$

# The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

**Average treatment effect (ATE):**

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

# The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

**Average treatment effect (ATE):**

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

This is the average of the individual treatment effects

# The Rubin causal model: average treatment effect

We can derive three parameters of interest from the definition of a treatment effect:

**Average treatment effect (ATE):**

$$E[\delta_i] = E[Y_i^1 - Y_i^0] = E[Y_i^1] - E[Y_i^0]$$

This is the average of the individual treatment effects

It is also unknowable

# Rubin: average treatment on the treated

**Average treatment on the treated (ATT):**

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

# Rubin: average treatment on the treated

**Average treatment on the treated (ATT):**

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the  $i$  in the treated group

# Rubin: average treatment on the treated

**Average treatment on the treated (ATT):**

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the  $i$  in the treated group

It is also unknowable since we never observe  $E[Y_i^0 | D_i = 1]$

# Rubin: average treatment on the treated

**Average treatment on the treated (ATT):**

$$E[\delta_i | D_i = 1] = E[Y_i^1 - Y_i^0 | D_i = 1] = E[Y_i^1 | D_i = 1] - E[Y_i^0 | D_i = 1]$$

This is the average of the individual treatment effects only for the  $i$  in the treated group

It is also unknowable since we never observe  $E[Y_i^0 | D_i = 1]$

If the treatment effect differs across  $i$  then  $ATE \neq ATT$

# Rubin: average treatment on the untreated

**Average treatment on the untreated (ATU):**

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

# Rubin: average treatment on the untreated

**Average treatment on the untreated (ATU):**

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the  $i$  in the untreated group

# Rubin: average treatment on the untreated

**Average treatment on the untreated (ATU):**

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the  $i$  in the untreated group

It is also unknowable since we never observe  $E[Y_i^1 | D_i = 0]$

# Rubin: average treatment on the untreated

**Average treatment on the untreated (ATU):**

$$E[\delta_i | D_i = 0] = E[Y_i^1 - Y_i^0 | D_i = 0] = E[Y_i^1 | D_i = 0] - E[Y_i^0 | D_i = 0]$$

This is the average of the individual treatment effects only for the  $i$  in the untreated group

It is also unknowable since we never observe  $E[Y_i^1 | D_i = 0]$

If the treatment effect differs across  $i$  then  $ATE \neq ATU$

# Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

# Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment  $D$  is whether a state has a conservation policy

# Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment  $D$  is whether a state has a conservation policy

The outcome  $Y$  is the number of bird species in that state

# Hands on: understanding treatment effects

We've got our definitions, now lets be a bit more clear as to what we are doing with some empirical examples putting our new R tools to work

Suppose treatment  $D$  is whether a state has a conservation policy

The outcome  $Y$  is the number of bird species in that state

We want to understand the causal effect of conservation policy on the number of species

# Hands on: understanding treatment effects

Here's our dataset of both potential outcomes:

```
cons_df # data frame of conservation treatment, outcomes
```

```
## # A tibble: 8 x 4
##   state    Y1     Y0 delta
##   <int> <dbl> <dbl> <dbl>
## 1     1     4     7    -3
## 2     2     7     9    -2
## 3     3     7     0     7
## 4     4    10     1     9
## 5     5     7     7     0
## 6     6     6     0     6
## 7     7     9     3     6
## 8     8    13     4     9
```

Calculate the average treatment effect

# Hands on: understanding treatment effects

ATE =  $E[\delta_i]$  which we can compute with `dplyr::summarise`:

```
cons_df %>%  
  dplyr::summarise(mean(delta))
```

```
## # A tibble: 1 x 1  
##   `mean(delta)`  
##       <dbl>  
## 1           4
```

The average treatment effect is 4: a conservation policy increases the number of bird species in a state by 4

# Hands on: understanding treatment effects

Notice that not all states benefit from conservation policies, and it even backfires in one state

The ATE is just the average over all the different treatment effects

```
cons_df
```

```
## # A tibble: 8 x 4
##   state    Y1     Y0  delta
##   <int> <dbl> <dbl> <dbl>
## 1     1     4     7    -3
## 2     2     7     9    -2
## 3     3     7     0     7
## 4     4    10     1     9
## 5     5     7     7     0
## 6     6     6     0     6
## 7     7     9     3     6
## 8     8    13     4     9
```

# Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

# Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

The policymaker assigns treatment and then observes the actual outcome according to  $Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$

# Hands on: understanding treatment effects

Suppose we have a perfect policymaker who knows each state's potential outcomes and can perfectly decide whether each state should have a conservation policy

The policymaker assigns treatment and then observes the actual outcome according to  $Y_i = D_i Y_i^1 + (1 - D_i) Y_i^0$

What does the dataset look like for the observed outcomes?

# Hands on: understanding treatment effects

```
observed_df ← cons_df %>%  
  mutate(  
    Y = ifelse(delta > 0, Y1, Y0),  
    D = as.numeric(delta > 0)  
  ) %>%  
  select(state, Y, D)  
observed_df
```

```
## # A tibble: 8 x 3  
##   state     Y     D  
##   <int> <dbl> <dbl>  
## 1     1     7     0  
## 2     2     9     0  
## 3     3     7     1  
## 4     4    10     1  
## 5     5     7     0  
## 6     6     6     1  
## 7     7     9     1  
## 8     8    13     1
```

# Hands on: estimating ATEs

Given the **observed** data, what if we tried to *estimate* the ATE by comparing mean outcomes of treated ( $D_i = 1$ ) vs untreated units ( $D_i = 0$ )

# Hands on: estimating ATEs

Given the **observed** data, what if we tried to *estimate* the ATE by comparing mean outcomes of treated ( $D_i = 1$ ) vs untreated units ( $D_i = 0$ )

This is the simple difference in mean outcomes (SDO):

$$\begin{aligned} SDO &= E[Y^1|D = 1] - E[Y^0|D = 0] \\ &= \frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_U} \sum_{i=1}^{N_U} (y_i | d_i = 0) \end{aligned}$$

where  $N_T$  is the number of treated units and  $N_U$  is the number of untreated units

# Hands on: estimating ATEs

We can compute the SDO using `dplyr::summarise` in conjunction with `dplyr::group_by` on treatment status  $D$ :

```
observed_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))
```

```
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    7.67  
## 2     1     9
```

The SDO is  $9 - 7.67 = 1.33 < 4$ !

# Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

# Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:<sup>1</sup>

$$\text{SDO} = \text{ATE} +$$

See Mixtape pages 89-91 for the derivation.

# Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:<sup>1</sup>

$$SDO = ATE + \text{Selection Bias} +$$

See Mixtape pages 89-91 for the derivation.

# Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:<sup>1</sup>

$$\text{SDO} = \text{ATE} + \text{Selection Bias} + \text{Heterogeneous Treatment Effect Bias}$$

See Mixtape pages 89-91 for the derivation.

# Hands on: estimating ATEs

We know the ATE is 4, why is the SDO giving us a much smaller estimate of the effect of the conservation policy?

Because the SDO is actually composed of three pieces, only one of which is the ATE:<sup>1</sup>

$$\text{SDO} = \text{ATE} + \text{Selection Bias} + \text{Heterogeneous Treatment Effect Bias}$$

What are these mathematically and intuitively?

See Mixtape pages 89-91 for the derivation.

# Hands on: decomposing the SDO

$$\underbrace{E[Y^1|D=1] - E[Y^0|D=0]}_{\text{SDO}} = \underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) + \frac{1}{N_U} \sum_{i=1}^{N_U} (y_i | d_i = 0)}_{\text{SDO}}$$
$$= \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$
$$+ \underbrace{E[Y^0|D=1] - E[Y^0|D=0]}_{\text{Selection bias}}$$
$$+ \underbrace{(1 - \pi)(ATT - ATU)}_{\text{Het. Treat. Eff. Bias}}$$

where  $\pi$  is the share of treated units

# Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

# Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

# Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

We **know** ATE = 4 so the last two terms are why SDO < ATE

# Hands on: decomposing the SDO

The LHS we know sums to 1.33 so the RHS must as well

The first term is the ATE, what we actually want to estimate

We **know** ATE = 4 so the last two terms are why SDO < ATE

Let's work through these two in more detail

# Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

# Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

# Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

The problem is we don't observe  $E[Y^0|D = 1]$ : what the treated group ( $D = 1$ ) would have looked like without treatment ( $Y^0$ )

# Hands on: Selection bias

The second term is **selection bias**:

$$E[Y^0|D = 1] - E[Y^0|D = 0]$$

It is the inherent differences between the two groups if they did not *actually* get the conservation policy

how are the two groups different at baseline?

The problem is we don't observe  $E[Y^0|D = 1]$ : what the treated group ( $D = 1$ ) would have looked like without treatment ( $Y^0$ )

Calculate this with the `cons_df` data frame with both potential outcomes

# Hands on: Selection bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(Y0)) # Difference in potential control outcomes across the two groups
```

```
## # A tibble: 2 x 2
##       D `mean(Y0)`
##   <dbl>     <dbl>
## 1     0      7.67
## 2     1      1.6
```

Selection bias is thus  $1.6 - 7.67 = -6.07$

# Hands on: het. treat. effect bias

The third term is the bias from heterogeneous treatment effects across groups:

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

It is the difference in the effect of the conservation policy across the two groups multiplied by the share that did not get a conservation policy

# Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

# Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group

# Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group
- Scale the size of this difference by fraction of control units

# Hands on: het. treat. effect bias

$$\underbrace{(1 - \pi)}_{\text{Share w/o policy}} \times \underbrace{(ATT - ATU)}_{\text{Diff. in treat. effect}}$$

What's the intuition for this?

- How much more of an effect did the policy have on the units that *happened* to get treatment versus the effect the policy would have had on the units that *happened* to be in the control group
- Scale the size of this difference by fraction of control units

Calculate this with the `cons_df` data frame with both potential outcomes

# Hands on: het. treat. effect bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(delta), n()) # Difference in potential control and treatment outcomes across the two groups
```

```
## # A tibble: 2 x 3
##       D `mean(delta)` `n()`
##   <dbl>     <dbl>   <int>
## 1     0      -1.67     3
## 2     1       7.4      5
```

Heterogeneous treatment effect bias is thus:  $(1 - 5/(5+3)) * (7.4 - (-1.67)) = 3.40$

# Hands on: het. treat. effect bias

```
cons_df %>%
  mutate(
    D = as.numeric(delta > 0)
  ) %>%
  group_by(D) %>%
  summarise(mean(delta), n()) # Difference in potential control and treatment outcomes across the two groups
```

```
## # A tibble: 2 x 3
##       D `mean(delta)` `n()`
##   <dbl>     <dbl>   <int>
## 1     0      -1.67     3
## 2     1       7.4      5
```

Heterogeneous treatment effect bias is thus:  $(1 - 5/(5+3)) * (7.4 - (-1.67)) = 3.40$

In total we have: **SDO (1.33) = ATE (4) + SB (-6.07) + HTEB (3.40)**

# Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

# Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

# Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

1. **Selection bias:** The units in the treatment group different from the control group in the absence of treatment

# Recap: treatment effect estimates

Taking a simple difference in means of outcomes between treatment and control groups **does** contain what we want

Unfortunately it is **confounded** by two forms of bias:

1. **Selection bias:** The units in the treatment group different from the control group in the absence of treatment
2. **Heterogeneous treatment effect bias:** The units in the treatment group respond to treatment differently than units in the control group

# Bias examples

What are some examples of these forms of bias?

# Bias examples

What are some examples of these forms of bias?

**Selection bias:**

# Bias examples

What are some examples of these forms of bias?

**Selection bias:** If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

# Bias examples

What are some examples of these forms of bias?

**Selection bias:** If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

**HTEB:**

# Bias examples

What are some examples of these forms of bias?

**Selection bias:** If we select certain groups into treatment, for example, if we pass conservation policy in states with little biodiversity

- This may lead to a negative bias on conservation policy impacts, even with a positive ATE

**HTEB:** If we select units into treatment based on expected response, for example, if we pass policy in states where birds are **very** sensitive to conservation

- This may lead to an overestimate of the size of the treatment effect

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid
4. Seeing how this all works (hands on) in modern environmental economics

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Empirical economics

A huge chunk of economics is trying circumvent these forms of bias<sup>1</sup>

The rest of class is largely dedicated to:

1. Understanding what tools we have to correctly estimate the ATE
2. Understanding what assumptions are required for these tools to be valid
3. Understanding when these ATE estimates are valid
4. Seeing how this all works (hands on) in modern environmental economics

This should be a less-technical complement to Brian's class

<sup>1</sup>We're more heavily concerned with selection bias than HTEB. There are also other biases to worry about that we will get to later.

# Recovering the ATE

There are many different ways to recover the ATE

# Recovering the ATE

There are many different ways to recover the ATE

We will cover some subset of:

1. Randomized control trials
2. Regression discontinuity
3. Difference-in-differences
4. Cross-sectional regressions
5. Two way fixed effects

# Recovering the ATE

There are many different ways to recover the ATE

We will cover some subset of:

1. Randomized control trials
2. Regression discontinuity
3. Difference-in-differences
4. Cross-sectional regressions
5. Two way fixed effects

All of these approaches have pluses and minuses

# Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

# Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

# Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

Selection bias is because treatment is correlated with unit characteristics in the absence of treatment

# Understanding the source of bias

What is the fundamental problem leading to selection bias and HTEB?

Both stem from treatment being correlated with characteristics of the observational units

Selection bias is because treatment is correlated with unit characteristics in the absence of treatment

HTEB is because treatment is correlated with the size of units' responses to treatment

# A potential solution

What's a simple way to deal with this?

# A potential solution

What's a simple way to deal with this?

**RANDOMIZATION**

# A potential solution

What's a simple way to deal with this?

## RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

# A potential solution

What's a simple way to deal with this?

## RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

Randomization drives selection bias and HTEB to zero

# A potential solution

What's a simple way to deal with this?

## RANDOMIZATION

If we randomize treatment across observational units then SDO = ATE

Randomization drives selection bias and HTEB to zero

Let's see why

# Randomization

Randomization means that treatment will be **independent** of the potential outcomes

# Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

# Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

In our example we know independence is violated because we assigned the conservation policy to states that had  $Y^1 > Y^0$

# Randomization

Randomization means that treatment will be **independent** of the potential outcomes

$$Y^1, Y^0 \perp D$$

In our example we know independence is violated because we assigned the conservation policy to states that had  $Y^1 > Y^0$

What if we randomized conservation policy?

# Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

# Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

# Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

This means that:

$$E[Y^1|D = 1] = E[Y^1|D = 0] \quad E[Y^0|D = 1] = E[Y^0|D = 0]$$

# Randomization

Randomization of treatment / policy means that:

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

If we randomized treatment across units, then on average, the difference in potential outcomes across groups should be zero

This means that:

$$E[Y^1|D = 1] = E[Y^1|D = 0] \quad E[Y^0|D = 1] = E[Y^0|D = 0]$$

The second equation  $E[Y^0|D = 1] - E[Y^0|D = 0] = 0$  directly gives us that selection bias is zero with the SDO

# Randomization

What does randomization do to HTEB?

# Randomization

What does randomization do to HTEB?

$$ATT - ATU$$

$$= (E[Y^1|D = 1] - E[Y^0|D = 1]) - (E[Y^1|D = 0] - E[Y^0|D = 0])$$

$$= (E[Y^1|D = 1] - E[Y^1|D = 0]) - (E[Y^0|D = 1] - E[Y^0|D = 0])$$

$$= (0) - (0) = 0$$

# Randomization

What does randomization do to HTEB?

$$ATT - ATU$$

$$= (E[Y^1|D = 1] - E[Y^0|D = 1]) - (E[Y^1|D = 0] - E[Y^0|D = 0])$$

$$= (E[Y^1|D = 1] - E[Y^1|D = 0]) - (E[Y^0|D = 1] - E[Y^0|D = 0])$$

$$= (0) - (0) = 0$$

HTEB goes to zero!

# Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

# Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

# Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment

# Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment
2. Data on observable outcomes

# Randomization and SDO

Randomization means that SDO = ATE:

$$\underbrace{\frac{1}{N_T} \sum_{i=1}^{N_T} (y_i | d_i = 1) - \frac{1}{N_C} \sum_{i=1}^{N_C} (y_i | d_i = 0)}_{\text{SDO}} = \underbrace{E[Y^1] - E[Y^0]}_{\text{ATE}}$$

All we need to estimate the average treatment effect of a policy is:

1. Data on treatment assignment
2. Data on observable outcomes
3. The independence assumption:  $Y^1, Y^0 \perp D$

# Hands on: randomization and SDO

Let's see how this works in practice by re-constructing our dataset and then randomizing treatment

# Hands on: randomization and SDO

Let's see how this works in practice by re-constructing our dataset and then randomizing treatment

```
set.seed(12345)
ate <- 4 # average treatment effect
n_obs <- 100 # number of observations
cons_rand_df <- tibble(
  state = seq(1, n_obs), # state identifier
  Y0 = floor(runif(n_obs)*10)) %>% # control/untreated potential outcome
  mutate(
    D = as.numeric(runif(n()) > 0.5), # randomized treatment
    Y1 = Y0 + ate + round(rnorm(n()))), # generate treatment potential outcome
    Y = D*Y1 + (1-D)*Y0 # generate observed outcome
  ) %>%
  select(
    state, D, Y # keep only observable variables
  )
```

# Hands on: randomization and SDO

```
cons_rand_df # data frame of randomized treatment, observable outcome
```

```
## # A tibble: 100 x 3
##   state      D      Y
##   <int> <dbl> <dbl>
## 1     1      0      7
## 2     2      1     11
## 3     3      1     11
## 4     4      1     11
## 5     5      1      8
## 6     6      1      4
## 7     7      1      7
## 8     8      1     11
## 9     9      0      7
## 10    10     0      9
## # ... with 90 more rows
```

# Hands on: randomization and SDO

Now take the SDO:

```
cons_rand_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))
```

```
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    4.42  
## 2     1    8.82
```

The SDO is  $8.82 - 4.42 = 4.40$ , a very close estimate of the ATE!

# Hands on: randomization and SDO

Now take the SDO:

```
cons_rand_df %>%  
  dplyr::group_by(D) %>%  
  dplyr::summarise(meanY = mean(Y))
```

```
## # A tibble: 2 x 2  
##       D   meanY  
##   <dbl> <dbl>  
## 1     0    4.42  
## 2     1    8.82
```

The SDO is  $8.82 - 4.42 = 4.40$ , a very close estimate of the ATE!

As  $n_{\text{obs}} \rightarrow \infty$ , we will have that  $SDO \rightarrow ATE$

# Randomization: what it does and doesn't do

What does independence imply?

# Randomization: what it does and doesn't do

What does independence imply?

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

The two groups have the same potential outcomes, *on average*

# Randomization: what it does and doesn't do

What does independence imply?

$$E[Y^1|D = 1] - E[Y^1|D = 0] = 0 \quad E[Y^0|D = 1] - E[Y^0|D = 0] = 0$$

The two groups have the same potential outcomes, *on average*

What does independence **not** imply? That:

$$E[Y^1|D = 1] - E[Y^0|D = 0] = 0 \quad E[Y^1|D = 1] - E[Y^0|D = 1] = 0$$

It does not imply that the observed outcomes are the same across the two groups, nor does it imply that the two potential outcomes of a single group are the same

# Independence

Is independence is a reasonable assumption in **observational data?**

# Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

# Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

Independence is **unlikely** to hold in these scenarios

# Independence

Is independence is a reasonable assumption in **observational data**?

Observational data are just data that we collect from the real world

- e.g. observed state policy choices and outcomes, observed pollution levels and outcomes

Independence is **unlikely** to hold in these scenarios

So what can we do?

# Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

# Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables  $X$ , the potential outcomes are independent of treatment  $(Y^1, Y^0 \perp D)$

# Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables  $X$ , the potential outcomes are independent of treatment  $(Y^1, Y^0 \perp D)$

This is much weaker than (unconditional) independence: we only need independence to hold for units that share the same  $X$  values, not for all units

# Conditional independence

Often we will instead rely on **conditional independence**:

$$(Y^1, Y^0 \perp D) | X$$

After controlling for some other variables  $X$ , the potential outcomes are independent of treatment  $(Y^1, Y^0 \perp D)$

This is much weaker than (unconditional) independence: we only need independence to hold for units that share the same  $X$  values, not for all units

Many of the estimation tools used in economics rely on (variants of) the conditional independence assumption